

RAKEZ Case Study: Deploying and Monitoring a Lead Scoring Model

Role: Machine Learning Engineer

Duration: 5-7 days

Status: ■ Complete

Executive Summary

This document presents a comprehensive solution for deploying, monitoring, and maintaining RAKEZ's lead scoring model in a production environment. The solution addresses all requirements including deployment strategy, online testing, monitoring, automation, and retraining capabilities.

Key Deliverables:

- Production-ready deployment architecture
- Online testing strategy (A/B testing and shadow deployment)
- Comprehensive monitoring and alerting system
- Automated retraining pipeline
- Real-time monitoring dashboard
- Complete documentation and presentation

1. Deployment Strategy

Architecture Overview

The deployment architecture leverages **Databricks** for data processing and **MLflow** for model management, with a **FastAPI** REST API for real-time serving.

Key Components

1.1 Model Registry (MLflow)

Implementation: 03_api/fastapi_app.py, 02_notebooks/model_inference_databricks.py

- **Stages:** Production, Staging, Archived
- **Versioning:** Automatic version tracking with metadata
- **Auditability:** Complete experiment tracking and model lineage
- **Rollback:** One-click reversion to previous versions

```
# Model loading from registry
model_uri = f"models://{MODEL_NAME}/{STAGE}"
model = mlflow.sklearn.load_model(model_uri)
```

1.2 REST API (FastAPI)

Implementation: 03_api/fastapi_app.py

- **Endpoint:** /score-lead for real-time scoring
- **Features:**

- Input validation (Pydantic models)
- Shadow model support
- Request logging
- Error handling
- Health checks

1.3 Batch Inference (Databricks)

Implementation: `02_notebooks/model_inference_databricks.py`

- Scheduled daily batch jobs
- Processes new leads from Delta Lake
- Updates CRM via Delta tables
- Handles feature engineering consistently

Deployment Methods

Canary Deployment

- Phase 1: 10% traffic (1 day)
- Phase 2: 50% traffic (2 days)
- Phase 3: 100% traffic
- Automatic rollback on failure

Shadow Deployment

- Parallel evaluation without affecting production
- Silent comparison of predictions
- Real-world performance data collection

Frameworks & Tools

- ■ **MLflow**: Model versioning and registry
- ■ **FastAPI**: REST API framework
- ■ **Databricks**: Data processing platform
- ■ **Delta Lake**: Data storage and versioning

2. Online Testing Approach

Strategy Overview

We implement a **two-phase testing approach**: Shadow deployment followed by A/B testing.

Phase 1: Shadow Deployment

Duration: 1-2 weeks

Implementation: `03_api/fastapi_app.py` (shadow model support)

- New model runs in parallel with production
- All traffic routed to both models
- Predictions compared silently
- Zero risk to business operations

Metrics Tracked:

- Prediction distribution differences
- Model performance (AUC, Precision, Recall)
- Business KPIs (conversion rate, revenue)

Phase 2: A/B Testing

Traffic Split: 90% production, 10% new model

Success Criteria:

- **Technical:** +2% AUC improvement OR
- **Business:** +5% conversion rate OR +5% revenue per lead
- **Operational:** No increase in error rate, latency maintained

Decision Process:

1. Monitor for statistical significance
2. Compare business metrics
3. Gather sales team feedback
4. Gradual rollout if successful

Safety Measures

- **Automatic Rollback:** Triggered if error rate > 2% or latency degrades > 50%
- **Real-time Monitoring:** Continuous tracking during testing
- **Approval Gates:** Manual review before full promotion
- **Gradual Rollout:** 10% → 50% → 100% traffic

3. Monitoring Plan

3.1 Data Drift Detection

Implementation: `02_notebooks/drift_detection.py`

Metrics:

PSI (Population Stability Index)

- **Calculation:** Per-feature and overall
- **Thresholds:**

- PSI < 0.1: No significant change
- PSI 0.1-0.25: Moderate change (Warning)
- PSI > 0.25: Significant change (Critical)

KL Divergence

- Measures distribution shift
- Threshold: $KL > 0.1$ indicates drift

Statistical Tests

- **Kolmogorov-Smirnov**: Continuous features
- **Chi-square**: Categorical features
- P-value < 0.05 indicates significant drift

Monitoring Schedule:

- **Real-time**: Feature statistics
- **Hourly**: Distribution checks
- **Daily**: PSI calculation
- **Weekly**: Comprehensive drift report

3.2 Prediction Drift

Implementation: `02_notebooks/monitoring_metrics.py`

- Model performance degradation detection
- AUC, Precision, Recall tracking over time
- Calibration metrics (Brier score)

3.3 Latency & Throughput

Metrics:

- **Latency**: P50, P95, P99 percentiles
- Target: P95 $< 200\text{ms}$
- Warning: P95 $> 500\text{ms}$
- **Throughput**: Requests per second
- Target: $> 10 \text{ req/s}$
- **Error Rate**: HTTP errors
- Target: $< 1\%$
- Critical: $> 1\%$

3.4 Business Performance

Metrics:

- **Conversion Rate**: Leads \rightarrow Customers
- **Revenue per Lead**: Average revenue
- **Score Distribution**: Lead score buckets
- **Conversion by Score**: Performance by score category

3.5 Alerting System

Implementation: `01_architecture/monitoring_architecture.md`

Alert Levels:

- **Level 1 (Info)**: Logged to dashboard

- **Level 2 (Warning):** Slack notification (#ml-alerts)
- **Level 3 (Critical):** Email + Slack + PagerDuty

Alert Triggers:

- PSI > 0.5: Critical drift
- Latency P95 > 500ms: Critical
- Error rate > 1%: Critical
- Conversion rate drop > 10%: Warning

3.6 Sales Team Complaint Investigation

Scenario: "Model scores are no longer helping prioritize leads effectively"

Investigation Workflow (6 Steps):

Immediate Response (1 hour)

- Check model health metrics
- Verify API functionality
- Review recent model changes
- Check data quality issues

Data Analysis (4 hours)

- Analyze conversion rates by score bucket
- Compare current vs historical performance
- Check for data drift (PSI, KL divergence)
- Review feature distributions

Model Performance Review (24 hours)

- Evaluate model metrics (AUC, Precision, Recall)
- Compare production vs shadow model
- Review calibration plots
- Check for concept drift

Root Cause Analysis

- **If Data Drift:** Investigate data source changes
- **If Concept Drift:** Business environment may have changed
- **If Model Issue:** Review training data and features
- **If Integration Issue:** Check CRM sync and data pipeline

Resolution

- **Data Issue:** Fix data pipeline, retrain model
- **Model Issue:** Retrain with updated data/features
- **Business Change:** Update model to reflect new patterns
- **Integration Issue:** Fix CRM sync mechanism

Communication & Prevention

- Document findings and resolution
- Update sales team with explanation
- Implement preventive measures
- Schedule follow-up review

Tools Used:

- Streamlit dashboard for metrics
 - MLflow UI for model comparison
 - Drift detection reports
 - Conversion rate analysis
-

4. Automation, Reproducibility & Retraining

4.1 Reproducibility

Implementation: MLflow + Delta Lake + Git

Components:

- **MLflow:** Experiment tracking, model versioning, parameter logging
- **Delta Lake:** Data snapshots and versioning
- **Git:** Code and configuration version control

Data Snapshots:

- Training data stored in Delta Lake with timestamps
- Feature engineering code versioned
- Model artifacts stored in MLflow

4.2 CI/CD Workflow

Implementation: 04_ci_cd/github_actions.yaml

Pipeline Stages:

Lint & Test

- Code formatting (Black)
- Linting (Flake8)
- Unit tests (Pytest)

Validate

- Notebook syntax validation
- Model validation checks

Deploy Staging

- Auto-deploy on develop branch
- Deploy notebooks to Databricks
- Trigger inference job

Deploy Production

- Manual approval required
- Canary deployment (10% → 50% → 100%)
- Monitor metrics
- Automatic rollback on failure

Model Registry Update

- Promote model from Staging to Production
- Update model tags and metadata

4.3 Retraining Strategy

Implementation: 02_notebooks/retraining_pipeline.py

Triggers:

Drift-Triggered

- Automatic when PSI > 0.25
- Significant distribution shift detected

Scheduled

- Weekly for first 3 months
- Monthly thereafter

Manual

- Admin-initiated for special cases

Retraining Process:

Data Collection

- Latest 6 months of labeled data
- Minimum 10,000 records required
- Time-based train/test split

Model Training

- Hyperparameter optimization (Optuna, 50-100 trials)
- Time series cross-validation
- XGBoost/LightGBM models

Model Evaluation

- Must outperform production model
- Minimum improvement: +2% AUC OR +5% business KPI
- Shadow testing for 1-2 weeks

Model Promotion

- Register to MLflow Staging
- Manual review and approval
- Canary deployment
- Promote to Production

Rollback Mechanism:

- **Automatic:** Error rate > 2%, Latency > 50% degradation, Business KPI drop > 10%
- **Manual:** One-click rollback to previous version
- **History:** Maintains model registry history

5. Bonus Features

5.1 Monitoring Dashboard

Implementation: `05_dashboard/streamlit_dashboard.py`

Features:

- Real-time metrics visualization
- Drift detection charts (PSI by feature)
- Performance trends (latency, throughput)
- Business metrics (conversion rates by score bucket)
- Alert viewer
- Works in demo mode (mock data) and production mode

Tabs:

1. **Overview:** System metrics and performance trends
2. **Drift Detection:** PSI and distribution monitoring
3. **Performance:** Latency and throughput metrics
4. **Business Metrics:** Conversion rates and score distributions
5. **Alerts:** System alerts and notifications

5.2 CRM Integration

Implementation: 06_docs/detailed_readme.md (CRM Integration Plan)

3-Phase Approach:

Phase 1: Batch Integration

- Scheduled job updates CRM scores daily
- Delta table as intermediary
- Low latency requirement

Phase 2: Real-time Integration

- FastAPI sends webhook on prediction
- CRM receives score immediately
- Higher latency requirement

Phase 3: API Integration

- CRM calls FastAPI directly
- On-demand scoring
- Highest latency requirement

Integration Methods:

- Webhook integration
- API integration
- Database integration (Delta tables)

5.3 Feedback Loop

Implementation: 06_docs/presentation_slides.md (Slide 10)

- Sales team complaint investigation workflow
- Root cause analysis process
- Preventive measures implementation
- Regular stakeholder communication

Technical Implementation

Code Structure

```
rakez-lead-scoring-deployment/  
■■■ 01_architecture/      # Architecture diagrams  
■■■ 02_notebooks/         # Databricks notebooks  
■■■ 03_api/               # FastAPI application  
■■■ 04_ci_cd/             # CI/CD pipeline  
■■■ 05_dashboard/        # Monitoring dashboard  
■■■ 06_docs/              # Documentation
```

Key Technologies

- **ML:** XGBoost, LightGBM, scikit-learn
- **Platform:** Databricks, MLflow

- **API:** FastAPI, Uvicorn
- **Dashboard:** Streamlit, Plotly
- **CI/CD:** GitHub Actions
- **Data:** Delta Lake, Spark

Conclusion

This solution provides a **complete, production-ready system** for deploying and monitoring RAKEZ's lead scoring model. All assessment requirements have been met:

- **Deployment Strategy:** Complete with MLflow, FastAPI, and Databricks
- **Online Testing:** Shadow deployment and A/B testing implemented
- **Monitoring Plan:** Comprehensive drift detection, performance metrics, and alerting
- **Automation & Retraining:** CI/CD pipeline and automated retraining
- **Bonus Features:** Working dashboard, CRM integration, feedback loop

The solution is **enterprise-grade, well-documented**, and **ready for production deployment**.

Appendix

Files and Locations

- **Architecture Diagrams:** 01_architecture/
- **Production Code:** 02_notebooks/, 03_api/
- **CI/CD Pipeline:** 04_ci_cd/github_actions.yaml
- **Dashboard:** 05_dashboard/streamlit_dashboard.py
- **Documentation:** 06_docs/

Quick Start

```
# Setup
python setup.py

# Start services
python start.py

# Access
# - API: http://localhost:8000/docs
# - Dashboard: http://localhost:8501
```

End of Case Study